

Assn2PartA-GrpNo43

October 6, 2023

0.1 # Pattern Recognition And Machine Learning

0.2 Assignment 2 - Part A

Group No: 43 Group Members: - Ajul Thomas (u3253992) - Hamad Rasheed (u3224704)

0.3 # Heart Disease Prediction using Pattern Recognition and Machine Learning

0.4 Data set selection and description of dataset and features.

0.5 Description

Our project focuses on developing a Predictive Risk Model for Heart Disease (PRMHD) using Pattern Recognition and Machine Learning (PRML). The model aims to analyze a range of medical and lifestyle factors to predict the likelihood of an individual developing heart disease.

0.6 Motivation

Heart disease is a leading cause of death worldwide. Early diagnosis can lead to effective treatment, but traditional diagnostic methods are often slow and expensive. A computational model can provide quick, accurate, and cost-effective risk assessment.

0.7 Dataset

The Cleveland Heart Disease dataset from the UCI Machine Learning Repository was selected for this project. Although the database contains a total of 76 attributes, our focus is on a subset of 14 key attributes, as these are the ones most commonly cited in published research. Specifically, we are utilizing only the Cleveland database for this endeavor. The initial steps involve data cleaning to address any missing values, followed by Exploratory Data Analysis (EDA) to gain insights into dataset.

0.8 Dataset Description:

- 1) **age** - age of the individual in years
- 2) **sex** - 1 = Male , 2 = Female
- 3) **cp** - chest pain type
 - 1 - typical angina
 - 2 - atypical angina
 - 3 - non-anginal pain
 - 4 - asymptomatic

- 4) **trtbps** - resting blood pressure (in mm Hg on admission to the hospital)
- 5) **chol** - serum cholesterols in mg/dl
- 6) **fbs** - fasting blood sugar > 120 mg/dl
 - 1 - true
 - 0 - false
- 7) **restecg** - resting electrocardiographic results
 - 0 - normal
 - 1 - having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV)
 - 2 - showing probable or definite left ventricular hypertrophy by Estes' criteria
- 8) **thalach** - maximum heart rate achieved
- 9) **exng**: exercise induced angina
 - 1 - Yes
 - 0 - No
- 10) **oldpeak** - ST depression induced by exercise relative to rest
- 11) **slp** - the slope of the peak exercise ST segmen
 - 1 - Upsloping
 - 2 - Flat
 - 3 - Downsloping
- 12) **ca** - number of major vessels (0–3) colored by fluoroscopy
- 13) **thall** - 3 = normal; 6 = fixed defect; 7 = reversible defect

0.8.1 Import required libraries

```
[1]: import pandas as pd
import numpy as np
```

0.8.2 Import the Heart Disease Dataset

```
[2]: data = pd.read_csv('./data/heart.csv')
# data = pd.read_csv('./data/processed.cleveland.data')
data.head()
```

```
[2]:
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	ca	\
0	63	1	3	145	233	1	0	150	0	2.3	0	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	0	
3	56	1	1	120	236	0	1	178	0	0.8	2	0	
4	57	0	0	120	354	0	1	163	1	0.6	2	0	

	thall	output
0	1	1
1	2	1
2	2	1
3	2	1
4	2	1

0.9 ## Exploratory Data Analysis

EDA involves understanding the dataset's structure and basic statistics.

```
[3]: data.shape
```

```
[3]: (303, 14)
```

```
[4]: data.columns
```

```
[4]: Index(['age', 'sex', 'cp', 'trtbps', 'chol', 'fbs', 'restecg', 'thalachh',  
         'exng', 'oldpeak', 'slp', 'ca', 'thall', 'output'],  
         dtype='object')
```

```
[5]: data.dtypes
```

```
[5]: age           int64  
sex           int64  
cp           int64  
trtbps       int64  
chol         int64  
fbs          int64  
restecg      int64  
thalachh     int64  
exng         int64  
oldpeak      float64  
slp          int64  
ca           int64  
thall        int64  
output       int64  
dtype: object
```

```
[6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 303 entries, 0 to 302  
Data columns (total 14 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   age         303 non-null   int64  
1   sex         303 non-null   int64  
2   cp          303 non-null   int64  
3   trtbps      303 non-null   int64  
4   chol        303 non-null   int64  
5   fbs         303 non-null   int64  
6   restecg     303 non-null   int64  
7   thalachh    303 non-null   int64  
8   exng        303 non-null   int64  
9   oldpeak     303 non-null   float64
```

```

10 slp      303 non-null    int64
11 ca       303 non-null    int64
12 thall    303 non-null    int64
13 output   303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB

```

```
[7]: data.describe().T
```

```

[7]:
      count      mean      std   min   25%   50%   75%   max
age      303.0   54.366337   9.082101   29.0   47.5   55.0   61.0   77.0
sex      303.0    0.683168   0.466011    0.0    0.0    1.0    1.0    1.0
cp       303.0    0.966997   1.032052    0.0    0.0    1.0    2.0    3.0
trtbps   303.0  131.623762  17.538143   94.0  120.0  130.0  140.0  200.0
chol     303.0  246.264026  51.830751  126.0  211.0  240.0  274.5  564.0
fbs      303.0    0.148515   0.356198    0.0    0.0    0.0    0.0    1.0
restecg  303.0    0.528053   0.525860    0.0    0.0    1.0    1.0    2.0
thalachh 303.0  149.646865  22.905161   71.0  133.5  153.0  166.0  202.0
exng     303.0    0.326733   0.469794    0.0    0.0    0.0    1.0    1.0
oldpeak  303.0    1.039604   1.161075    0.0    0.0    0.8    1.6    6.2
slp      303.0    1.399340   0.616226    0.0    1.0    1.0    2.0    2.0
ca       303.0    0.729373   1.022606    0.0    0.0    0.0    1.0    4.0
thall    303.0    2.313531   0.612277    0.0    2.0    2.0    3.0    3.0
output   303.0    0.544554   0.498835    0.0    0.0    1.0    1.0    1.0

```

0.10 ## Data Cleaning

Checks and addresses if any missing values are there. If any will address these issues.

```
[8]: data.nunique()
```

```

[8]: age      41
     sex       2
     cp        4
     trtbps    49
     chol     152
     fbs       2
     restecg   3
     thalachh  91
     exng      2
     oldpeak   40
     slp       3
     ca        5
     thall     4
     output    2
dtype: int64

```

```
[9]: data['ca'].unique()
```

```
[9]: array([0, 2, 1, 3, 4], dtype=int64)
```

```
[10]: data.ca.value_counts()
```

```
[10]: ca
0     175
1      65
2      38
3      20
4       5
Name: count, dtype: int64
```

```
[11]: data.duplicated().sum()
```

```
[11]: 1
```

```
[12]: data.isnull().sum()
```

```
[12]: age          0
sex            0
cp            0
trtbps        0
chol          0
fbs           0
restecg       0
thalachh      0
exng          0
oldpeak       0
slp           0
ca            0
thall         0
output        0
dtype: int64
```

0.11 ## Data Visualization

Create visualizations to gain insights into the data distribution, relationships, and patterns. Uses libraries like Matplotlib and Seaborn.

```
[13]: import matplotlib.pyplot as plt
import seaborn as sns
```

changing the data for better visualization and plotting

```
[14]: df = data.copy(deep=True)
df['output'] = df.output.replace({1: "Disease", 0: "No_disease"})
df['sex'] = df.sex.replace({1: "Male", 0: "Female"})
df['cp'] = df.cp.replace({0: "typical_angina",
                        1: "atypical_angina",
```

```

                2:"non-anginal pain",
                3: "asymtomatic"})
df['exng'] = df.exng.replace({1: "Yes", 0: "No"})
df['fbs'] = df.fbs.replace({1: "True", 0: "False"})
df['slp'] = df.slp.replace({0: "upsloping", 1: "flat", 2:"downsloping"})
df['thall'] = df.thall.replace({1: "fixed_defect", 2: "reversable_defect", 3:
↪ "normal"})

```

```
[15]: df
```

```

[15]:    age    sex          cp  trtbps   chol    fbs  restecg  thalachh  \
0     63  Male    asymptomatic    145   233   True        0        150
1     37  Male  non-anginal pain    130   250  False        1        187
2     41  Female  atypical_angina    130   204  False        0        172
3     56  Male    atypical_angina    120   236  False        1        178
4     57  Female    typical_angina    120   354  False        1        163
..    ...    ...          ...    ...    ...    ...    ...    ...
298   57  Female    typical_angina    140   241  False        1        123
299   45  Male    asymptomatic    110   264  False        1        132
300   68  Male    typical_angina    144   193   True        1        141
301   57  Male    typical_angina    130   131  False        1        115
302   57  Female  atypical_angina    130   236  False        0        174

```

```

    exng  oldpeak      slp  ca          thall    output
0     No     2.3    upsloping  0      fixed_defect  Disease
1     No     3.5    upsloping  0  reversable_defect  Disease
2     No     1.4  downsloping  0  reversable_defect  Disease
3     No     0.8  downsloping  0  reversable_defect  Disease
4     Yes     0.6  downsloping  0  reversable_defect  Disease
..    ...    ...          ...  ..    ...    ...
298  Yes     0.2      flat  0          normal  No_disease
299   No     1.2      flat  0          normal  No_disease
300   No     3.4      flat  2          normal  No_disease
301  Yes     1.2      flat  1          normal  No_disease
302   No     0.0      flat  1  reversable_defect  No_disease

```

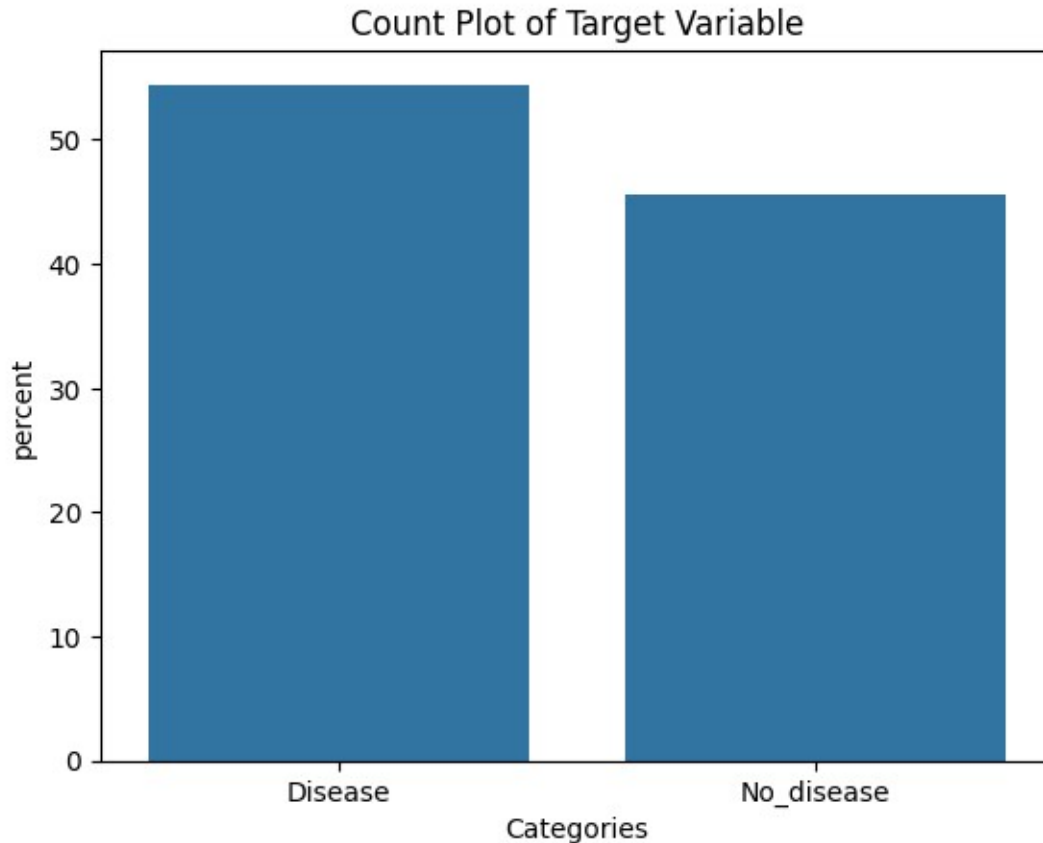
```
[303 rows x 14 columns]
```

0.11.1 Target Variable Distribution

```

[16]: sns.countplot(df, x='output', stat="percent")
plt.title('Count Plot of Target Variable')
plt.xlabel('Categories')
plt.show()

```



Observation:

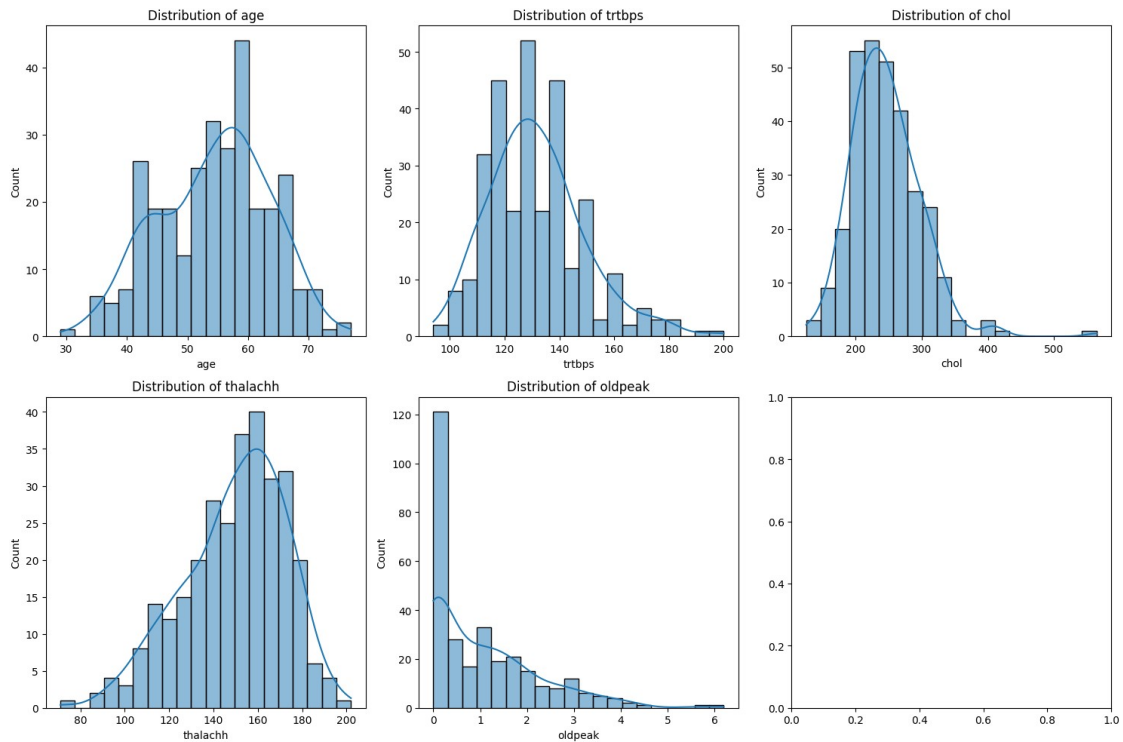
- Our initial graph categorizes individuals based on the presence or absence of heart disease. In the dataset, just over 50% of participants have heart disease, while approximately 45% do not. This balanced distribution provides a solid foundation for more in-depth analysis.

0.11.2 Distribution of Numerical Variables

```
[17]: numerical_features = ['age', 'trtbps', 'chol', 'thalachh', 'oldpeak']
fig, axes = plt.subplots(nrows=2, ncols=3, figsize=(15, 10))
fig.subplots_adjust(hspace=0.5)
for i, feature in enumerate(numerical_features):
    row, col = i // 3, i % 3
    ax = axes[row, col]
    sns.histplot(data[feature], bins=20, kde=True, ax=ax)
    ax.set_title(f'Distribution of {feature}')
    ax.set_xlabel(feature)
    ax.set_ylabel('Count')

# Adjust layout
```

```
plt.tight_layout()
plt.show()
```



```
[18]: # Assuming 'target' is the name of the column representing your target variable
target_variable = 'output'

fig, axes = plt.subplots(nrows=3, ncols=2, figsize=(10, 12))
fig.subplots_adjust(hspace=0.5)

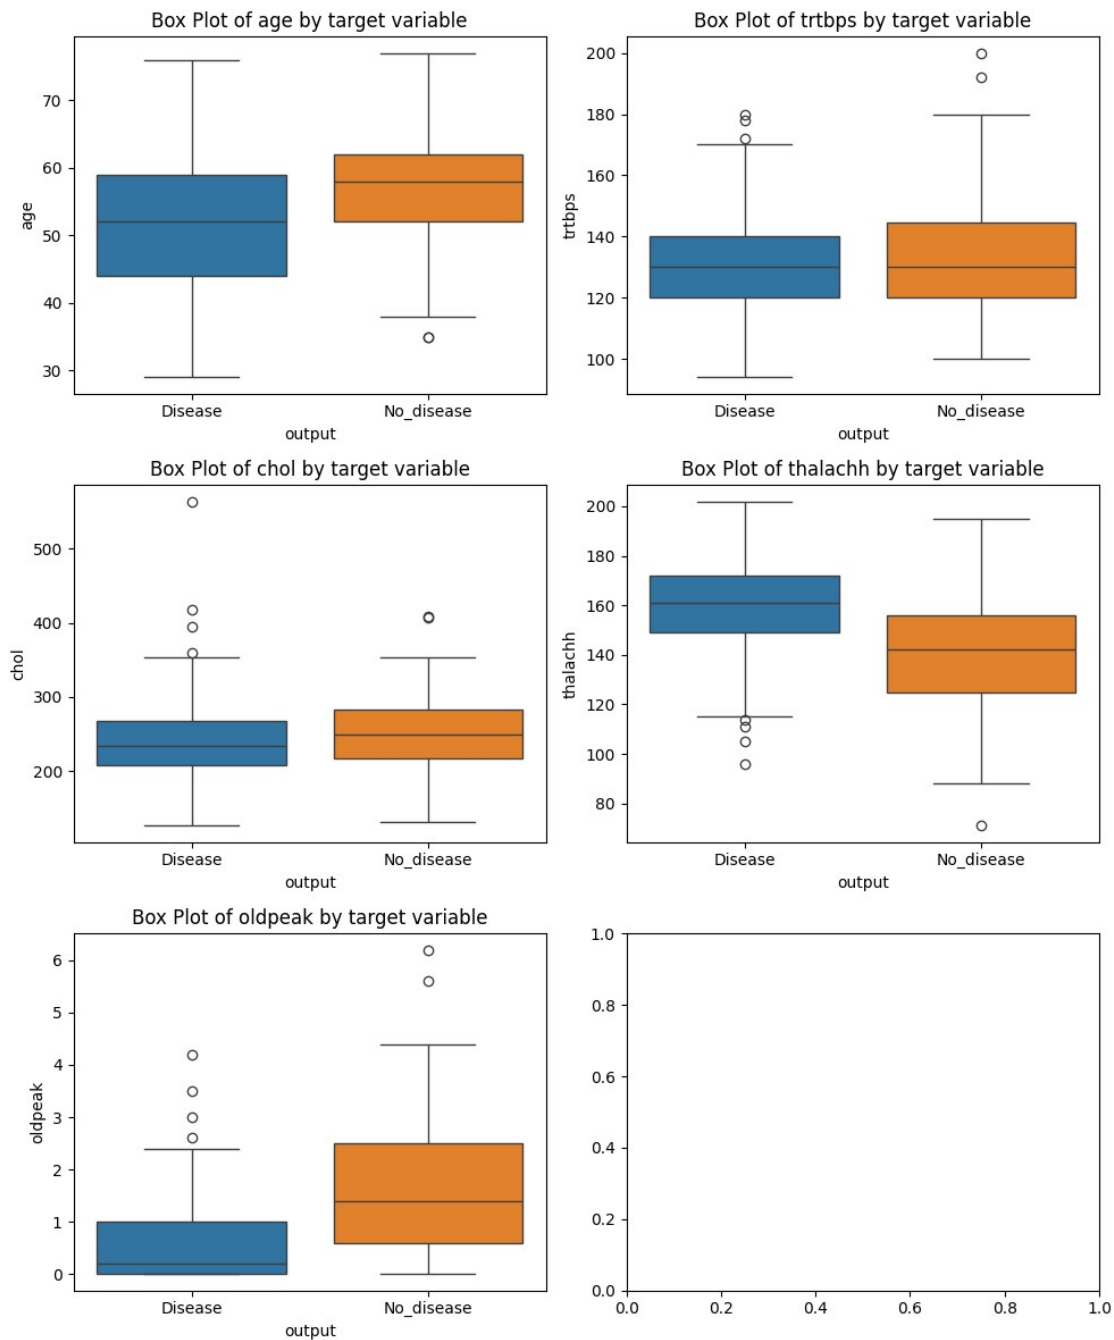
colors = ["blue", "green"]

# Loop through each attribute and plot a box plot for each target variable
for i, attribute in enumerate(numerical_features):
    row, col = i//2, i % 2
    ax = axes[row, col]
    # plt.figure(figsize=(10, 6)) # Set the figure size
    sns.boxplot(x=target_variable, y=attribute, hue=target_variable, data=df,
    ↪ax=ax)
    ax.set_title(f'Box Plot of {attribute} by target variable')
    ax.set_xlabel(target_variable)
    ax.set_ylabel(attribute)

# Adjust layout
```



```
plt.tight_layout()
plt.show()
```



Observations: As we examine the distributions of numerical variables such as age, resting heart rate (trtbps), cholesterol (chol), maximum heart rate achieved (thalachh), and exercise-induced ST depression (oldpeak).

The dataset reveals that the average age for those with heart disease is lower compared to those without it. Most individuals in the dataset are aged between 50 and 70, following a normal distribution.

In terms of cholesterol levels, there's little variation between those with and without heart disease, although some outliers are present. The majority of individuals have cholesterol levels ranging between 200 and 300, adhering to a normal distribution.

When considering maximum heart rate achieved, people with heart disease generally have higher rates than those without. Several outliers exist, but the majority have heart rates between 150 and 175.

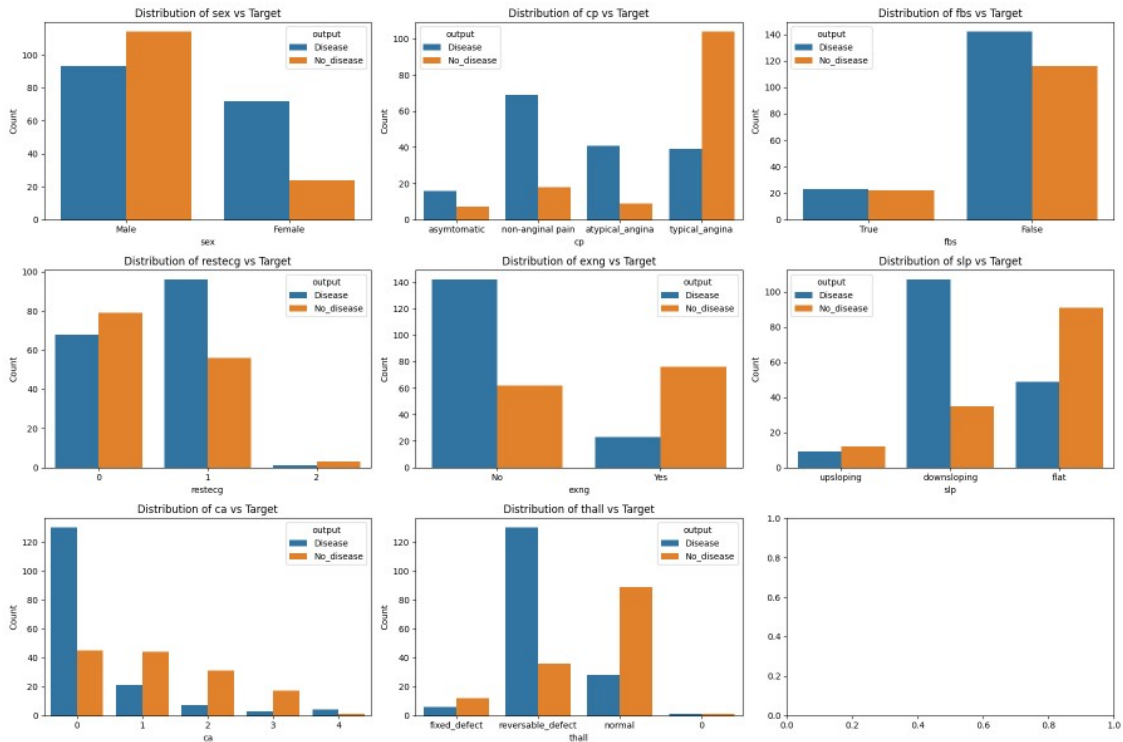
Lastly, the average oldpeak level for those with heart disease is lower than for those without. Despite some outliers, the distribution for oldpeak is right-skewed, with most individuals registering a value of zero.

0.11.3 Distribution of Categorical Variables

```
[19]: # Create bar plots for categorical features
categorical_features = ['sex', 'cp', 'fbs', 'restecg', 'exng', 'slp', 'ca', 'u
    ↪ 'thall']
fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(18, 12))
fig.subplots_adjust(hspace=0.5)

for i, feature in enumerate(categorical_features):
    row, col = i // 3, i % 3
    ax = axes[row, col]
    sns.countplot(x=df[feature], hue=df['output'], ax=ax)
    ax.set_title(f'Distribution of {feature} vs Target')
    ax.set_xlabel(feature)
    ax.set_ylabel('Count')

# Adjust layout
plt.tight_layout()
plt.show()
```




0.11.4 Gender distribution according to target variable

```
[20]: fig, ax = plt.subplots(figsize=(8,5))
name = df['sex']
ax = sns.countplot(x='sex', hue='output', data=df, palette='Set1')
ax.set_title("Sex Distribution according to Target", fontsize = 13, weight = 'bold')
ax.set_xticklabels (name, rotation = 0)

totals = []
for i in ax.patches:
    totals.append(i.get_height())
total = sum(totals)
for i in ax.patches:
    ax.text(i.get_x()+.05, i.get_height()-15,
            str(round((i.get_height()/total)*100, 2))+'%', fontsize=14,
            color='white', weight = 'bold')

plt.tight_layout()
```

C:\Users\ajult\AppData\Local\Temp\ipykernel_19948\3742478429.py:5: UserWarning: FixedFormatter should only be used together with FixedLocator
 ax.set_xticklabels (name, rotation = 0)



Chest pain distribution according to target variable

```
[21]: df.cp.value_counts()
```

```
[21]: cp
      typical_angina      143
non-anginal pain      87
      atypical_angina      50
      asymptomatic       23
Name: count, dtype: int64
```

```
[22]: fig, ax = plt.subplots(figsize=(10,5))
name = df.cp.unique()
ax = sns.countplot(x='cp', hue='output', data=df, palette='Set1')
ax.set_title("Chest Pain Distribution according to Target", fontsize = 13,
             weight = 'bold')
ax.set_xticklabels (name, rotation = 0)


totals = []
for i in ax.patches:
    totals.append(i.get_height())
total = sum(totals)
for i in ax.patches:
    ax.text(i.get_x()+.03, i.get_height()-5,
            str(round((i.get_height()/total)*100, 2))+'%', fontsize=14,
```

```

        color='white', weight = 'bold')
plt.tight_layout()

```

C:\Users\ajult\AppData\Local\Temp\ipykernel_19948\815883350.py:5: UserWarning: FixedFormatter should only be used together with FixedLocator
 ax.set_xticklabels (name, rotation = 0)



Fasting blood sugar distribution according to target variable

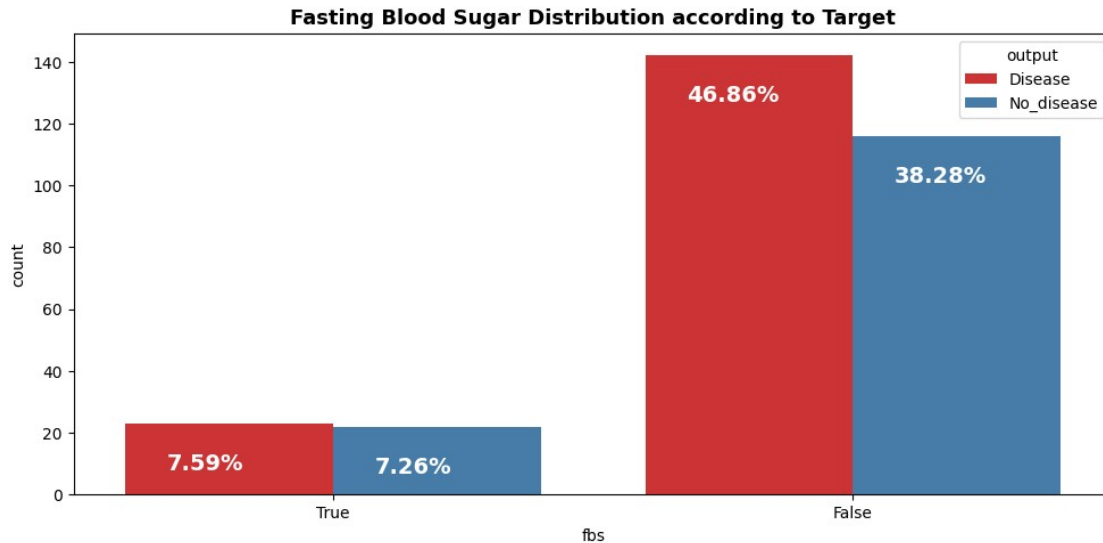
```

[23]: fig, ax = plt.subplots(figsize=(10,5))
name = df['fbs']
ax = sns.countplot(x='fbs', hue='output', data=df, palette='Set1')
ax.set_title("Fasting Blood Sugar Distribution according to Target", fontsize = 13, weight = 'bold')
ax.set_xticklabels (name, rotation = 0)

totals = []
for i in ax.patches:
    totals.append(i.get_height())
total = sum(totals)
for i in ax.patches:
    ax.text(i.get_x()+.08, i.get_height()-15,
            str(round((i.get_height()/total)*100, 2))+'%', fontsize=14,
            color='white', weight = 'bold')
plt.tight_layout()

```

C:\Users\ajult\AppData\Local\Temp\ipykernel_19948\2214451207.py:5: UserWarning: FixedFormatter should only be used together with FixedLocator
 ax.set_xticklabels (name, rotation = 0)



Observations on distribution of Categorical Variables vs Target: Sex vs. Target: The data indicates that a higher number of women have heart disease compared to those who don't, while the opposite is true for men. Men make up 68.3% of the study population.

Chest Pain (cp) vs. Target: Among the four levels of chest pain, individuals at level 2 are more prone to heart disease. Conversely, those at level 0 are less likely to have heart disease and make up 47.2% of the dataset.

Fasting Blood Sugar (fbs) vs. Target: Individuals with an fbs under 120 are more susceptible to heart disease and constitute 85.1% of the dataset.

Resting ECG (restecg) vs. Target: Those with a restecg result of 1 are more likely to have heart disease compared to those with a result of 0. The majority have results categorized as 0 or 1.

Exercise-Induced Angina (exang) vs. Target: Individuals without exercise-induced angina are more likely to have heart disease. This group represents 67.3% of the study population.

Slope of Peak Exercise ST Segment vs. Target: Those with a downslope are more susceptible to heart disease. Most individuals display either a flat or downslope.

Number of Major Vessels Colored by Fluoroscopy (CA) vs. Target: Participants with zero major vessels colored are more prone to heart disease, making up 57.8% of the dataset.

Thallium Stress Result (thal) vs. Target: Individuals with a thal value of 2 are more likely to have heart disease, and they constitute 54.8% of the study population.

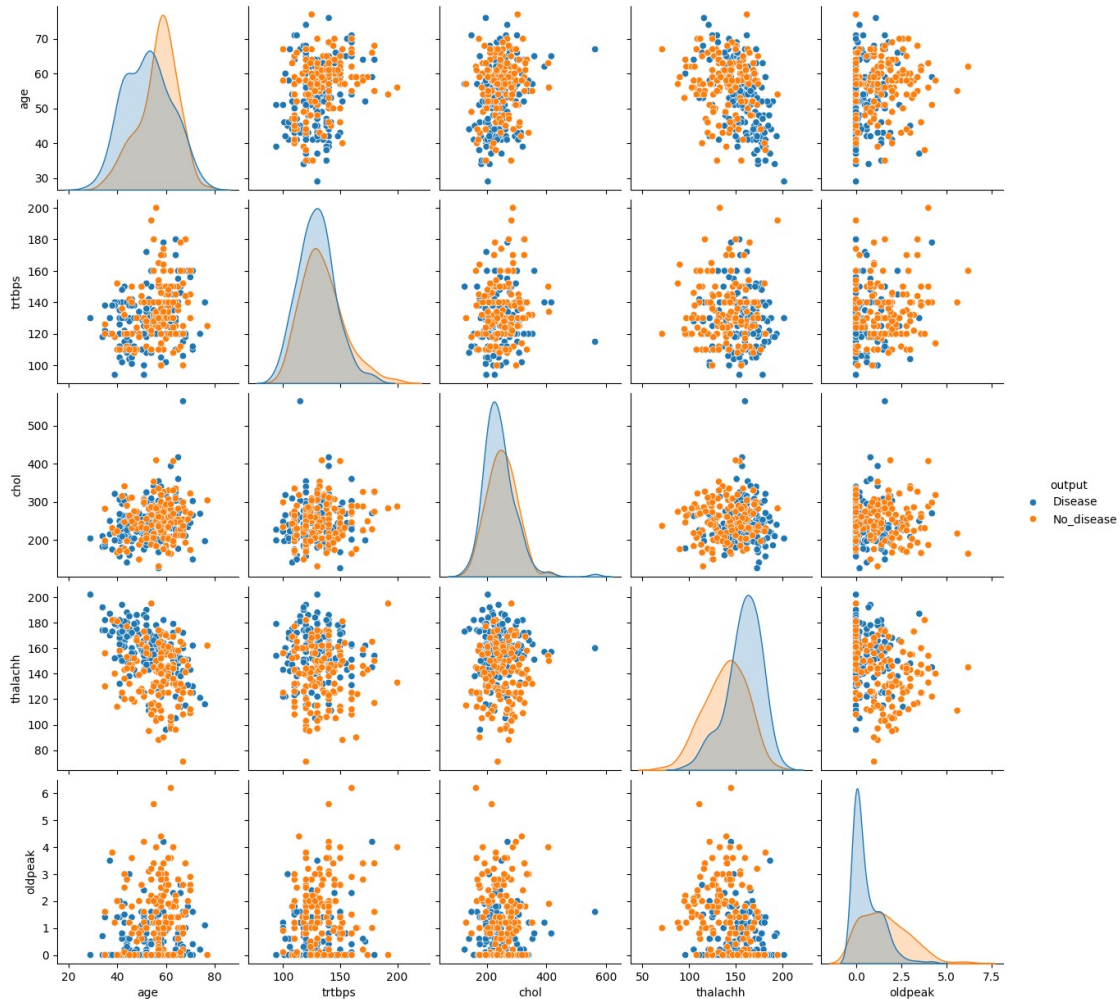
0.11.5 Visualize the distribution of continuous variable across target variable

```
[24]: sns.pairplot(df[numerical_features + ['output']], hue='output')
```

C:\Users\ajult\AppData\Roaming\Python\Python311\site-packages\seaborn\axisgrid.py:123: UserWarning: The figure layout has changed to

```
tight
self._figure.tight_layout(*args, **kwargs)
```

[24]: <seaborn.axisgrid.PairGrid at 0x248000ffc10>



Observation: A pair plot comparing heart disease to non-heart disease across numerical variables offers a comprehensive overview of the dataset, highlighting patterns and correlations among different metrics for both categories.

Age: The average age for individuals with heart disease is lower than for those without, suggesting age could inversely correlate with risk. The age distribution is mostly normal, centered around 50 to 70 years.

Resting Heart Rate (trtbps) & Cholesterol (chol): These variables show little variation between the heart disease and non-heart disease groups, indicating they may not be strong predictors. Outliers in these variables warrant further investigation.

Maximum Heart Rate Achieved (thalachh): Individuals with heart disease generally achieve higher

maximum heart rates, which could be an important variable for predictive modeling.


ST Depression Induced by Exercise Relative to Rest (oldpeak): A lower average oldpeak is observed among individuals with heart disease, potentially pointing to different stress responses between the two groups.

The pair plot serves as a valuable tool for visualizing interactions among these variables within the context of heart disease and non-heart disease categories

0.12 # Correlation Matrix

```
[25]: # Calculate the correlation matrix
correlation_matrix = data.corr()

# Visualize the correlation matrix as a heatmap
plt.figure(figsize=(18, 12))
sns.heatmap(correlation_matrix, annot=True, cmap='YlGnBu', linewidths=.5)
plt.title('Correlation Matrix Heatmap')
plt.show()
```



Observations: Lowest Correlation: Fasting Blood Sugar (fbs) and Cholesterol (chol) show the lowest correlation with the target variable. This aligns with earlier observations that these variables

exhibited little variation between individuals with and without heart disease, suggesting they may not be strong predictors.

General Correlations: Most other variables are correlated with each other and with the target variable. For instance, age has an inverse correlation with the likelihood of having heart disease, while maximum heart rate achieved (thalachh) tends to be higher in individuals with heart disease.

The correlation matrix can serve as a statistical foundation for more in-depth analysis, helping to identify key variables that could be central to predictive modeling for heart disease.

0.13 Selecting PRML Algorithms

Selecting appropriate machine learning algorithms for a specific dataset involves considering factors like the nature of the data, the problem type, and the desired outcomes. In the case of heart disease dataset, which is a binary classification problem (predicting the presence or absence of heart disease), below mentioned algorithms can be suitable:

1. **Logistic Regression:**

- **Applicability:** Binary classification problems.
- **Reasoning:** Logistic Regression is a simple and interpretable algorithm that can serve as a baseline model. It works well when the relationship between features and the target variable is approximately linear. Given your dataset has both numerical and categorical features, logistic regression can handle both types effectively.

2. **Random Forest:**

- **Applicability:** Classification problems, especially with structured data.
- **Reasoning:** Random Forest is an ensemble algorithm known for its ability to handle both numerical and categorical features. It is robust, provides feature importance scores, and often works well “out of the box.” Random Forests are also less prone to overfitting and can handle noisy data.

3. **Support Vector Machines (SVM):**

- **Applicability:** Binary classification problems.
- **Reasoning:** SVM is effective for binary classification tasks, even when the data is not linearly separable. It works well when there is a clear margin of separation between classes. SVM can handle both numerical features. s specific goals.

Reason for choosing the above algorithms

- 1) **Data Type:** Heart Disease dataset contains a mix of numerical and categorical features, which makes it suitable for algorithms like Random Forest, Gradient Boosting, and Logistic Regression that can handle both types of features effectively.
- 2) **Binary Classification:** Our task is binary classification (predicting the presence or absence of heart disease), making algorithms designed for classification tasks, like Logistic Regression, Random Forest, and Gradient Boosting, relevant.
- 3) **Complexity:** While Logistic Regression is simple and interpretable, Random Forest can capture more complex relationships in the data, which might be important for achieving high predictive accuracy.
- 4) **Ensemble Methods:** Random Forest and Gradient Boosting are both ensemble methods, which can help improve model performance by combining multiple weak learners.

0.14 Implementing Algorithms

```
[26]: # import required modules
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[27]: data.shape

# features
X = data.iloc[:, 0:-1]

# target variable
y = data.iloc[:, -1]
```

```
[28]: # split X and y into training and testing sets
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
                                                    random_state=16)
```

0.15 Spot Checking Algorithms with accuracy report

```
[29]: # Load libraries
from pandas import read_csv
from pandas.plotting import scatter_matrix
from matplotlib import pyplot
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC

seed = 7

# Spot-Check Algorithms
models = []
models.append(('LR', LogisticRegression(solver='liblinear', multi_class='ovr')))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
```

```

models.append(('CART', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))
models.append(('SVM', SVC(kernel='linear')))
# evaluate each model in turn
results = []
names = []
for name, model in models:
    kfold = KFold(n_splits=10, random_state=seed, shuffle=True)
    cv_results = cross_val_score(model, X_train, y_train, cv=kfold,
    ↪scoring='accuracy')
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)

# Compare Algorithms
fig = pyplot.figure()
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
pyplot.boxplot(results)
ax.set_xticklabels(names)
pyplot.show()


```

```

LR: 0.832016 (0.090642)
LDA: 0.831818 (0.073588)
KNN: 0.642095 (0.114810)
CART: 0.792490 (0.077942)
NB: 0.805731 (0.106107)
SVM: 0.840909 (0.062291)

```

Algorithm Comparison



0.15.1 Logistic Regression

```
[30]: # import the class
from sklearn.linear_model import LogisticRegression

# instantiate the model (using the default parameters)
logreg = LogisticRegression(random_state=16, max_iter=1000)

# fit the model with data
logreg.fit(X_train, y_train)

y_pred = logreg.predict(X_test)
```

```
[31]: # import the metrics class
from sklearn import metrics

cnf_matrix = metrics.confusion_matrix(y_test, y_pred)
cnf_matrix
```

```
[31]: array([[31, 10],
           [ 4, 31]], dtype=int64)
```

```
[32]: from sklearn.metrics import classification_report

report = classification_report(y_test, y_pred, target_names = ['Healthy',
↳ 'Heart Disease'])

print(report)
```

	precision	recall	f1-score	support
Healthy	0.89	0.76	0.82	41
Heart Disease	0.76	0.89	0.82	35
accuracy			0.82	76
macro avg	0.82	0.82	0.82	76
weighted avg	0.83	0.82	0.82	76

0.15.2 Support Vector Machines (SVM)

```
[33]: #Import svm model
from sklearn import svm

#Create a svm Classifier
clf = svm.SVC(kernel='linear') # Linear Kernel

#Train the model using the training sets
clf.fit(X_train, y_train)

#Predict the response for test dataset
y_pred_svm = clf.predict(X_test)
```

```
[34]: # import the metrics class
from sklearn import metrics

cnf_matrix = metrics.confusion_matrix(y_test, y_pred_svm)
cnf_matrix
```

```
[34]: array([[32,  9],
           [ 3, 32]], dtype=int64)
```

```
[35]: from sklearn.metrics import classification_report

report = classification_report(y_test, y_pred_svm, target_names = ['Healthy',
↳ 'Heart Disease'])
```

```
print(report)
```

	precision	recall	f1-score	support
Healthy	0.91	0.78	0.84	41
Heart Disease	0.78	0.91	0.84	35
accuracy			0.84	76
macro avg	0.85	0.85	0.84	76
weighted avg	0.85	0.84	0.84	76

0.15.3 Random Forest

```
[36]: from sklearn.ensemble import RandomForestClassifier
```

```
rf = RandomForestClassifier()  
rf.fit(X_train, y_train)
```

```
[36]: RandomForestClassifier()
```

```
[37]: y_pred_rf = rf.predict(X_test)
```

```
[38]: # import the metrics class  
from sklearn import metrics  
  
cnf_matrix = metrics.confusion_matrix(y_test, y_pred_rf)  
cnf_matrix
```

```
[38]: array([[33,  8],  
         [ 5, 30]], dtype=int64)
```

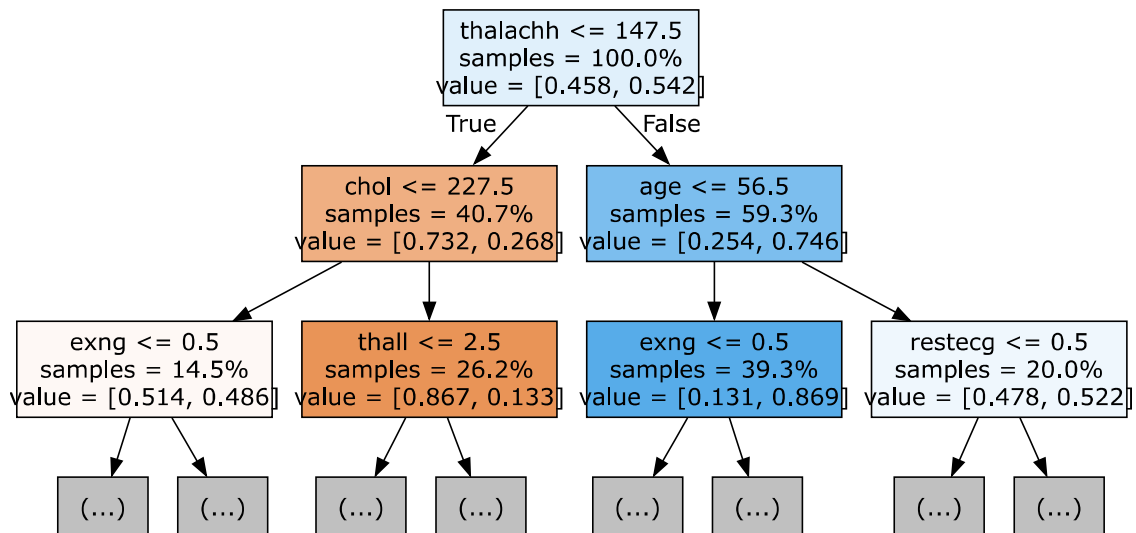
```
[39]: from sklearn.metrics import classification_report  
  
report = classification_report(y_test, y_pred_rf, target_names = ['Healthy',  
↪ 'Heart Disease'])  
  
print(report)
```

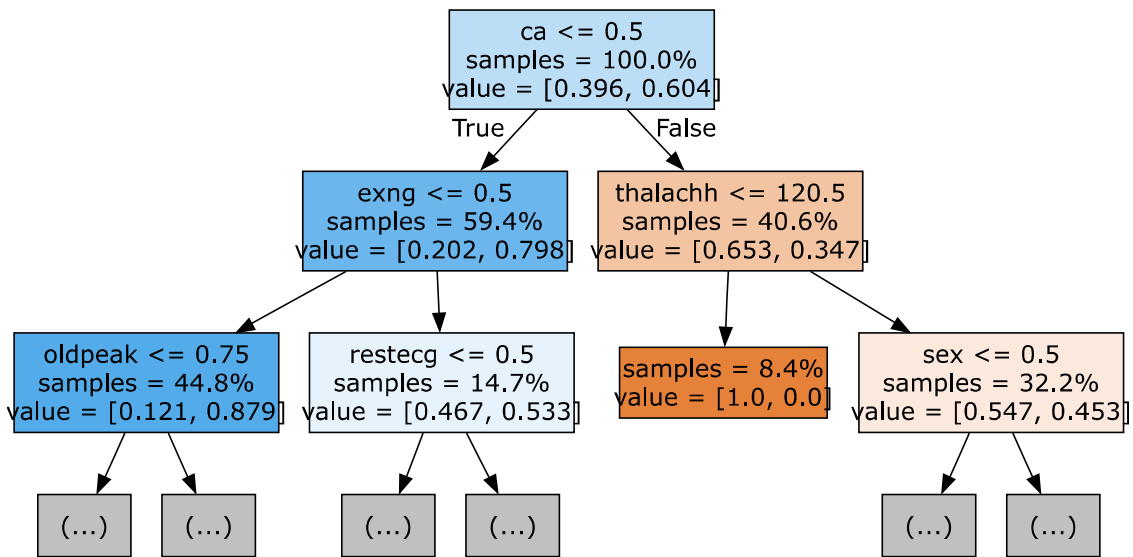
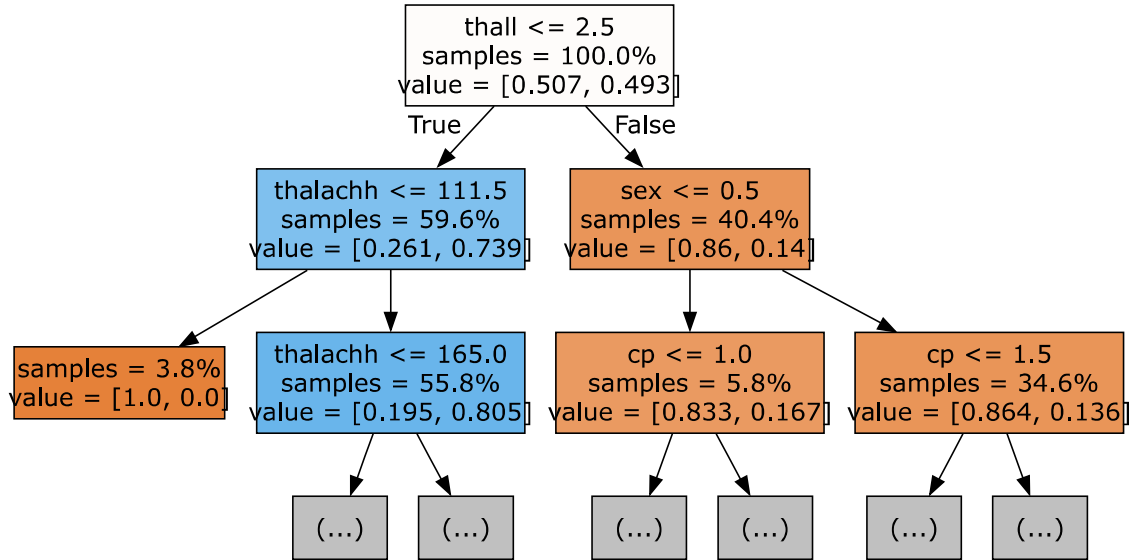
	precision	recall	f1-score	support
Healthy	0.87	0.80	0.84	41
Heart Disease	0.79	0.86	0.82	35
accuracy			0.83	76
macro avg	0.83	0.83	0.83	76
weighted avg	0.83	0.83	0.83	76

```
[40]: # Export the first three decision trees from the forest
from sklearn.tree import export_graphviz
from IPython.display import Image
import graphviz

for i in range(3):
    tree = rf.estimators_[i]
    dot_data = export_graphviz(tree,
                               feature_names=X_train.columns,
                               filled=True,
                               max_depth=2,
                               impurity=False,
                               proportion=True)

    graph = graphviz.Source(dot_data)
    display(graph)
```





0.16 ### References:

Author Unknown. (n.d.). Exploratory Data Analysis on Heart Disease UCI Data Set. Towards Data Science. Retrieved from <https://towardsdatascience.com/exploratory-data-analysis-on-heart-disease-uci-data-set-ae129e47b323>

Gaur, D. (n.d.). A Guide to any Classification Problem. Kaggle. Retrieved from <https://www.kaggle.com/code/durgancegaur/a-guide-to-any-classification-problem>

Rahmanpritom, R. (n.d.). Heart Attack Analysis & Prediction Dataset. Kaggle. Retrieved from <https://www.kaggle.com/datasets/rashikrahanpritom/heart-attack-analysis-prediction-dataset>